

Exascale Computing at the LHC : Narrative

Peter Elmer, Matthew Jones, Steven Ko,
Salvatore Rappoccio, Lukasz Zialek

April 30, 2013

1 Funding Strategy

This research proposal outlines the necessity to extend the computing capacity of the experiments at the Large Hadron Collider (LHC) in Geneva, Switzerland, to the exascale of high-throughput processing. This is an enormously challenging task, but a necessary one to ensure the long-term success of the LHC experiments.

Exascale computing (in this case, in high throughput) is an enormously growth-oriented area. Even during lean economic times, the US Federal Government is pledging to support this area of research, for instance in the Department of Energy Exascale Computing Initiative [1]. Indeed, the DOE Office of Science quotes :

The Exascale initiative will be significant and transformative for Department of Energy missions.

The current level of funding for the DOE EIC and related activities is \$21M. [2]. This “transformative” strategy is one that is envisioned to continue well into the future.

In addition to governmental programs such as above, private sector funding sources are also available, such as the Google Faculty Research Awards [3], which has an interest in exascale computing projects also.

Need more bla bla bla.

2 Project Organization

The principle investigators (PIs) of this proposal have a widely-varied and applicable skill set to accomplish the goals of extending LHC computing to the exascale.

- Salvatore Rappoccio has 15 years of experience programming in a high-energy physics environment, as well as other numerical software design for the private sector. He is an expert in critical areas of event reconstruction at CMS which can be optimized for multicore usage.
- Lukasz Ziarek
- Steven Ko
- Peter Elmer
- Matthew Jones

Add your bla bla bla above. .

3 Introduction

With the discovery of a new boson with mass around $m = 125$ GeV [4, 5] (henceforth referred to as the H), a new phase of particle physics has begun. The questions have shifted from the cause of the breaking of the electroweak symmetry, to the nature of that symmetry breaking. Two major questions arise. The first is the exact nature of the particle responsible for the electroweak symmetry breaking. The second is how a particle with a relatively low mass around $m = 125$ GeV can be responsible for electroweak symmetry breaking without extremely large fine tuning in nature, cancelling the large radiative corrections to its mass.

With 25 fb^{-1} of 7 and 8 TeV data delivered by the LHC in 2011-2012, and instantaneous luminosities reaching $7 \times 10^{33} \text{ cm}^{-2} \text{ s}^{-1}$, the processing time to reconstruct each event collected by CMS was approximately 20 seconds per event. However, as the instantaneous luminosity is increased, the computational time currently scales quadratically. As the upgraded LHC is expected to deliver $> 12 \times 10^{33} \text{ cm}^{-2} \text{ s}^{-1}$ in the upcoming run, the processing time per event is expected to reach several minutes per event as shown in Figure 1. Furthermore, in future runs of the LHC in the next 15 years, the luminosity is expected to reach as high as $> 1 \times 10^{34} \text{ cm}^{-2} \text{ s}^{-1}$, which would correspond (naively) to several hours of computational time per event! Clearly, it is necessary for the computing power to scale in order to compensate for this dramatic increase in CPU time with instantaneous luminosity.

However, with the expected end of the historic scaling of single-core processing capability [6] (**NEED MORE MATERIAL HERE**), it is imperative to utilize a parallel processing strategy in order to maintain the levels of computational speed of LHC data in the immediate future in experiments such as CMS.

As described in detail in Ref. [7], there are several opportunities to utilize a parallel strategy in order to improve CPU usage. In particular, one application that has many possible opportunities is in combining the decay products of hadronic “jets” of particles via a nearest-neighbor clustering algorithm (“jet clustering”). This is a well-established technique employed at many different particle physics experiments worldwide, and is implemented in a common software framework called **fastjet** [8]. The single-core optimization of the mathematical implementation of the nearest-neighbor (NN) algorithm chosen is outlined in Ref. [9], and depends on the number of “candidates” that are input to the jet clustering algorithm. The single-core opti-

mization yields $O(N^2)$ or $O(N \ln N)$ operational times. This is analogous to the “K-nearest neighbors algorithm” [10] (kNN).

However, even though the single-core computational strategy is sufficient for many applications, there are two key components of jet clustering that are still inefficient for future LHC data processing, which can be solved with appropriate development of a parallelization strategy. The first is the parallelization of the NN algorithm itself to make optimum use of the more advanced vectorization capabilities of modern multicore CPUs. The second is to use a “divide and conquer” strategy to reconstruct a single collision event in several disjoint sections of the CMS detector in parallel.

There is existing work and literature on the topic of the parallelization of the kNN algorithm, for instance, in Refs. [11, 12, 13], where improvements $O(100)$ in CPU performance are observed over standard algorithms. Since the proposed use case is very similar to the kNN algorithm, similar improvements to the processing time by parallelization strategies are expected. Furthermore, a critical piece of information in the jet algorithm is the “area” of each individual jet (in the NN-related metric). Currently, the procedure to estimate this area is to add a large number of infinitesimally small candidates called “ghosts” that are uniformly distributed over the total area. These are fed into the jet clustering algorithm, and the positions of these ghosts are used to estimate the size. The resolution of this area is inversely related to the number of ghosts per unit area, so in practice many ghosts ($O(1000)$) are added in the neighborhoods of the jets. This severely limits the computational speed with which these can be processed. Judicious usage of parallelization of the algorithm may be able to solve this problem, and drastically reduce the computational time needed to compute the area.

In addition to the usage of GPU and other parallelization strategies to estimate the jet area, it is also possible to parallelize the computation by dividing the event into disjoint sections, and computing the jet clustering in parallel over multiple cores. This optimization is factorized from the previous, in that it can be done with current single-core computational algorithms, but simply divides the problem into several smaller ones. There are both computational and physics-related challenges that need to be overcome in this strategy, and using a synergistic approach is absolutely critical.

Need more bla bla bla.




Figure 1: Event processing time versus instantaneous luminosity.

References

- [1] US Department of Energy Office of Science, (2010), The Opportunities and Challenges of Exascale Computing.
- [2] US Department of Energy Office of Science, (2013), ASCR Budget.
- [3] Google, (2013), Faculty Research Awards.
- [4] CMS Collaboration, S. Chatrchyan et al., Phys. Lett. B (2012), 1207.7235.
- [5] ATLAS Collaboration, G. Aad et al., Phys. Lett. B (2012), 1207.7214.
- [6] K. Bergman et al., Defense Advanced Research Projects Agency Information Processing Techniques Office (DARPA IPTO) (2008), Exascale computing study: Technology challenges in achieving exascale systems.
- [7] Google, (2013), Faculty Research Awards.
- [8] M. Cacciari, G.P. Salam and G. Soyez, Eur.Phys.J. C72 (2012) 1896, 1111.6097.
- [9] M. Cacciari and G.P. Salam, Phys.Lett. B641 (2006) 57, hep-ph/0512210.

- [10] T. Cover and P. Hart, Information Theory, IEEE Transactions on 13 (1967) 21.
- [11] V. Garcia et al., IEEE International Conference on Image Processing (ICIP), Hong Kong, China, 2010.
- [12] V. Garcia, E. Debreuve and M. Barlaud, CVPR Workshop on Computer Vision on GPU, Anchorage, Alaska, USA, 2008.
- [13] V. Garcia, Suivi d'objets d'intrt dans une squence d'images : des points saillants aux mesures statistiques, PhD thesis, Universit de Nice - Sophia Antipolis, Sophia Antipolis, France, 2008.